


Fashion MNIST Classification

Fashion-MNIST is a dataset of Zalando's article images—consisting of a **training set of 60,000** examples and a **test set of 10,000 examples**. Each example is a **28x28 grayscale** image, associated with a label from **10 classes**. We intend Fashion-MNIST to serve as a direct drop-in **replacement for the original MNIST** dataset for benchmarking machine learning algorithms. It shares the same image size and structure of training and testing splits.

Here's an example how the data looks (each class takes three-rows):



Step # 1 - Import Libraries

Lets import all the libraries we are going to require for this classification project. It is always good to put all the import statements at the beginning of the file.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sbn
        5 from sklearn.model_selection import train_test_split
        6 from sklearn.metrics import confusion_matrix, classification_report
        7 from keras.models import Sequential
        8 from keras.layers import Conv2D, MaxPooling2D, Dropout, Dense, Flatten
        9 from keras.optimizers import Adam
       10 from keras.callbacks import TensorBoard
       11 from keras.utils import to_categorical
```

Step # 2 - Load Data

Now lets use **pandas** library to read the train and test datasets in the respective csv files. We are going to use the **read_csv** function which reads a csv file and returns a pandas **DataFrame** object.

```
In [2]: 1 fashion_train_df = pd.read_csv('fashion-mnist_train.csv', sep=',')
        2 fashion_test_df = pd.read_csv('fashion-mnist_test.csv', sep=',')
```

Now that we have loaded the datasets, lets check some parameters about the datasets.

```
In [3]: 1 fashion_train_df.shape    # Shape of the dataset
```

```
Out[3]: (60000, 785)
```

```
In [4]: 1 fashion_train_df.columns    # Name of the columns of the DataSet.
```

```
Out[4]: Index(['label', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5', 'pixel6',
              'pixel7', 'pixel8', 'pixel9',
              ...,
              'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779', 'pixel780',
              'pixel781', 'pixel782', 'pixel783', 'pixel784'],
              dtype='object', length=785)
```

So we can see that the 1st column is the label or target value for each row.

Now Lets find out how many distinct lables we have.

```
In [5]: 1 print(set(fashion_train_df['label']))
```

```
{0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
```

So we have 10 different lables. from 0 to 9.

Now lets find out what is the min and max of values of in the other columns.

```
In [6]: 1 print([fashion_train_df.drop(labels='label', axis=1).min(axis=1).min(),
        2         fashion_train_df.drop(labels='label', axis=1).max(axis=1).max()])
```

```
[0, 255]
```

So we have 0 to 255 which is the color values for grayscale. 0 being white and 255 being black.

Now lets check some of the rows in tabular format

```
In [7]: 1 fashion_train_df.head()
```

```
Out[7]:
```

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pixel777	pixel778	pixel779
0	2	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
1	9	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
2	6	0	0	0	0	0	0	0	5	0	...	0	0	0	30	4
3	0	0	0	0	1	2	0	0	0	0	...	3	0	0	0	0
4	3	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0

5 rows × 785 columns



So every other thing of the test dataset are going to be the same as the train dataset except the shape.

```
In [8]: 1 fashion_test_df.shape
```

```
Out[8]: (10000, 785)
```

So here we have 10000 images instead of 60000 as in the train dataset.

Lets check first few rows.

```
In [9]: 1 fashion_test_df.head()
```

```
Out[9]:
```

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pixel777	pixel778	pixel779
0	0	0	0	0	0	0	0	0	9	8	...	103	87	56	0	0
1	1	0	0	0	0	0	0	0	0	0	...	34	0	0	0	0
2	2	0	0	0	0	0	0	14	53	99	...	0	0	0	0	6
3	2	0	0	0	0	0	0	0	0	0	...	137	126	140	0	13
4	3	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0

5 rows × 785 columns




Step # 3 - Visualization

Now that we have loaded the data and also got somewhat acquainted with it lets visualize the actual images. We are going to use **Matplotlib** library for this.

```

In [10]: 1 # Convert the dataframe to numpy array
2 training = np.asarray(fashion_train_df, dtype='float32')
3
4 # Lets show multiple images in a 15x15 grid
5 height = 10
6 width = 10
7
8 fig, axes = plt.subplots(nrows=width, ncols=height, figsize=(17,17))
9 axes = axes.ravel() # this flattens the 15x15 matrix into 225
10 n_train = len(training)
11
12 for i in range(0, height*width):
13     index = np.random.randint(0, n_train)
14     axes[i].imshow(training[index, 1:].reshape(28,28))
15     axes[i].set_title(int(training[index, 0]), fontsize=8)
16     axes[i].axis('off')
17
18 plt.subplots_adjust(hspace=0.5)

```



Step # 4 - Preprocess Data

Great! We have visualized the images. So now we can start preparing for creating our model. But before that we need to preprocess our data so that we can fit our model easily. Lets do that first.

Since we are dealing with image data and our task is to recognize and classify images our model should be a Convolutional Neural Network. For that our images should have atleast 3 dimensions (**height x width x color_channels**). But our images are flattened in one dimension, **784 pixel (28x28x1)** values per row. So we need to reshape the data into its original format.

```
In [11]: 1 # convert to numpy arrays and reshape
2 training = np.asarray(fashion_train_df, dtype='float32')
3 X_train = training[:, 1:].reshape([-1,28,28,1])
4 X_train = X_train/255 # Normalizing the data
5 y_train = training[:, 0]
6
7 testing = np.asarray(fashion_test_df, dtype='float32')
8 X_test = testing[:, 1:].reshape([-1,28,28,1])
9 X_test = X_test/255 # Normalizing the data
10 y_test = testing[:, 0]
```

Also we need to have three different sets of data for **training**, **validatin** and **testing**. We already have different sets for training and testing. So we are going to split the training dataset further into two sets and will use one set of training and the other for validation.

```
In [13]: 1 # Split training set into training and validation sets
X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.2, random_state=5) # TODO :
```

```
In [14]: 1 # Lets check the shape of all three datasets
2 print(X_train.shape, X_val.shape, X_test.shape)
3 print(y_train.shape, y_val.shape, y_test.shape)

(38400, 28, 28, 1) (9600, 28, 28, 1) (10000, 28, 28, 1)
(38400,) (9600,) (10000,)
```

Step # 5 - Create and Train the Model

Create the model

```
In [15]: 1 cnn_model = Sequential()
2 cnn_model.add(Conv2D(filters=64, kernel_size=(3,3), input_shape=(28,28,1), activation='relu'))
3 cnn_model.add(MaxPooling2D(pool_size = (2,2)))
4 cnn_model.add(Dropout(rate=0.3))
5 cnn_model.add(Flatten())
6 cnn_model.add(Dense(units=32, activation='relu'))
7 cnn_model.add(Dense(units=10, activation='sigmoid'))
```

compile the model

```
In [16]: 1 cnn_model.compile(optimizer=Adam(lr=0.001), loss='sparse_categorical_crossentropy', metrics=['acc'])
2         cnn_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d (MaxPooling2D)	(None, 13, 13, 64)	0
dropout (Dropout)	(None, 13, 13, 64)	0
flatten (Flatten)	(None, 10816)	0
dense (Dense)	(None, 32)	346144
dense_1 (Dense)	(None, 10)	330
Total params: 347,114		
Trainable params: 347,114		
Non-trainable params: 0		

C:\Users\admin\anaconda3\lib\site-packages\keras\optimizers\legacy\adam.py:117: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
 super().__init__(name, **kwargs)

Train the model

```
In [17]: 1 cnn_model.fit(x=X_train, y=y_train, batch_size=512, epochs=50, validation_data=(X_val, y_val))
```

loss: 0.2723 - val_accuracy: 0.9139
 Epoch 45/50
 75/75 [=====] - 20s 264ms/step - loss: 0.1146 - accuracy: 0.9589 - val_loss: 0.2765 - val_accuracy: 0.9145
 Epoch 46/50
 75/75 [=====] - 20s 272ms/step - loss: 0.1146 - accuracy: 0.9591 - val_loss: 0.2777 - val_accuracy: 0.9132
 Epoch 47/50
 75/75 [=====] - 20s 273ms/step - loss: 0.1159 - accuracy: 0.9570 - val_loss: 0.2761 - val_accuracy: 0.9128
 Epoch 48/50
 75/75 [=====] - 21s 278ms/step - loss: 0.1110 - accuracy: 0.9598 - val_loss: 0.2960 - val_accuracy: 0.9106
 Epoch 49/50
 75/75 [=====] - 22s 290ms/step - loss: 0.1104 - accuracy: 0.9595 - val_loss: 0.2875 - val_accuracy: 0.9137
 Epoch 50/50
 75/75 [=====] - 20s 272ms/step - loss: 0.1099 - accuracy: 0.9601 - val_loss: 0.2833 - val_accuracy: 0.9151

Step # 5 - Evaluate the Model

Get the accuracy of the model

```
In [30]: 1 eval_result = cnn_model.evaluate(X_test, y_test)
2         print("Accuracy : {:.3f}".format(eval_result[1]))
```


313/313 [=====] - 2s 6ms/step - loss: 0.2629 - accuracy: 0.9193
 Accuracy : 0.919

Visualize the model's predictions

```
In [44]: 1 y_pred = cnn_model.predict(x=X_test)
```

```
313/313 [=====] - 2s 5ms/step
```

```
In [20]: 1 height = 10
2 width = 10
3
4 fig, axes = plt.subplots(nrows=width, ncols=height, figsize=(20,20))
5 axes = axes.ravel()
6 for i in range(0, height*width):
7     index = np.random.randint(len(y_pred))
8     axes[i].imshow(X_test[index].reshape((28,28)))
9     axes[i].set_title("True Class : {:.0f}\nPrediction : {:.0f}".format(y_test[index],y_pred[index]))
10    axes[i].axis('off')
11 plt.subplots_adjust(hspace=0.8, wspace=0.5)
```



Plot Confusin Matrix